

CALEB KAN

☎ +44 0782 8210751 | ✉ calebkan1106@gmail.com | [in linkedin.com/in/caleb-kan](https://www.linkedin.com/in/caleb-kan) | github.com/caleb-kan | calebkan.com

EDUCATION

Imperial College London

London, United Kingdom

Degree: Bachelor of Engineering – BEng, Computing

Sep 2023 – Jun 2026

Grade: First-Class Honours (Predicted)

SKILLS

Languages: Native: English, Mandarin, Cantonese

Programming Languages: C, Haskell, HTML/CSS, Java, Kotlin, Prolog, Python, Scala, SQL

Databases: Firebase, Milvus, MongoDB, MySQL

Frameworks: Django, Flask, LangChain, Streamlit

Libraries: BeautifulSoup4, Matplotlib, NLTK, NumPy, OpenCV, Pandas, Playwright, PyTorch, Ray, SciPy, Selenium, SymPy

Developer Tools: CI/CD, Docker, GDB, Git, Google Colab, Hugging Face, JetBrains, Linux, Ollama, Valgrind, VS Code

EXPERIENCE

Software Engineer Intern

Jul 2025 – Sep 2025

HubSpot

London, United Kingdom

Software Engineer Intern

Aug 2024 – Oct 2024

Midas Advisory

London, United Kingdom

- Automated non-operating expense data processing from the top 15 U.S. banks using open-source LLMs, web scraping, data source APIs, and parallelisation with Ray, cutting processing time from hours to minutes and boosting accuracy to 97%.
- Implemented a Milvus vector database with high-dimensional embeddings using the BAAI/bge-m3 model. Built a scalable architecture for future ML model and data source integration.
- Developed a hybrid search system combining Milvus's RRFRanker and dual LLMs: a small LLM achieving 93% in data filtering accuracy and large LLM (Hermes-3-8B-GGUF) for structured schema-based reporting.

Research Intern

Jul 2023 – Jul 2023

Department of Computer Science, City University of Hong Kong

Hong Kong

- Achieved 92.5% accuracy in GDPR compliance detection across 400+ websites by developing a GDPR compliance checker extension utilising BeautifulSoup4, LangChain, and OpenAI's LLM for automated analysis of a website's privacy policy.
- Increased daily active users by 45%, improved user-reported understanding (from 3.2 to 4.7/5), and boosted successful data erasure requests by 150% (100+ in first month) through a one-click feature that automates sending requests to websites.

PROJECTS

WACC (Compiler Project) | Scala

Jan 2025 – Mar 2025

- Developed a WACC compiler frontend, building a lexer and parser for syntax analysis, creating an abstract syntax tree representation, implementing a symbol table for semantic checks, and designing descriptive error reporting to aid debugging.
- Implemented the backend using TAC intermediate representation with ARM32/AArch64 support, architecture-specific dependencies, code optimisations (constant propagation, folding, control flow analysis), and a standard math library.

PintOS (Operating Systems Project) | C

Oct 2024 – Dec 2024

- Enhanced OS kernel functionality by implementing timer-based thread synchronisation, advanced priority scheduling with priority donation, and the BSD scheduler, ensuring efficient multitasking and thread management.
- Developed a virtual memory subsystem, including paging, frame management with second-chance eviction, and supplementary page tables, enabling support for user programs, memory-mapped files, and stack growth.

ARMv8 Emulator, Assembler, and Visualiser | C

May 2024 – Jun 2024

- Engineered a cycle-accurate ARMv8 emulator and two-pass assembler, implementing precise register, memory, and instruction management including robust error handling in compliance with ARMv8 specifications.
- Developed a SDL2-based GUI for real-time emulation visualisation, featuring drag-and-drop assembly parsing and dynamic rendering of CPU states, registers, memory maps, and ALU.

AI Research Agent | Python

Aug 2023 – Sep 2023

- Developed a AI research agent using LangChain, OpenAI's LLM, Serper API, BeautifulSoup4, and X API for automated search, web scraping, cited summaries, and content posting. Integrated MongoDB Atlas for geolocation tracking.
- Implemented a Streamlit web app enabling anyone to use the AI research agent, integrating geospatial visualisation to map worldwide usage using geolocation data stored in MongoDB Atlas. Analysed engagement patterns of 70+ global users.

HONOURS AND AWARDS

1st Place – G-Research London Coding Challenge 2024

Issued by G-Research

- Achieved First Place among 40+ contestants in the G-Research London Coding Challenge, a 3-hour programming contest where participants compete to achieve the highest score on a given problem.