CALEB KAN

☎ +44 0782 8210751 | ✉ calebkan1106@gmail.com | in linkedin.com/in/caleb-kan | ⓞ github.com/caleb-kan | ⚇ calebkan.com

EDUCATION

Imperial College London   London, United Kingdom
Degree: Master of Engineering – MEng, Computing (Artificial Intelligence and Machine Learning)   Sep 2023 – Jun 2027
Grade: First-Class Honours (Year 1, Year 2)

SKILLS

Languages: **Native**: English, Mandarin, Cantonese
Programming Languages: C, Haskell, HTML/CSS, Java, Kotlin, Prolog, Python, Scala, SQL, TypeScript
Databases: Firebase, Milvus, MongoDB, MySQL, Qdrant, Snowflake, Supabase
Frameworks & Tech: Django, Expo, FastAPI, Flask, Jest, JUnit, Kafka, LangChain, React Native, ScalaTest, Streamlit
Developer Tools: CI/CD, CMake, Docker, GDB, Git, JetBrains, Jupyter, Linux, Node.js, Valgrind, Vercel, VS Code

EXPERIENCE

Software Engineer Intern   Jul 2025 – Sep 2025
HubSpot   London, United Kingdom
- Improved AI Agent tool reliability by 97.5% (error rate 40% to 1%) by implementing resilient domain parsing and URL normalisation; ran statistical timeout tuning for the generate image tool, achieving 0% errors across AI Agent executions.
- Architected an end-to-end agent refinement platform using Kafka, Snowflake, Qdrant (semantic search), and Python REST APIs to capture user refinements and apply LLM analysis for personalised, continuously improving agent behavior.
- Enhanced platform robustness with JSON Schema validation and strict type checking, blocking hallucinated parameters before execution to preserve data integrity and prevent failures across HubSpot's AI Agent ecosystem.

Software Engineer Intern   Aug 2024 – Oct 2024
Midas Advisory   London, United Kingdom
- Automated non-operating expense data processing from the top 15 U.S. banks using open-source LLMs, web scraping, data-source APIs, and Ray parallelisation, cutting processing time from hours to minutes and boosting accuracy to 97%.
- Implemented a Milvus vector-database with BAAI/bge-m3 embeddings and RRFRanker hybrid search, pairing a 93% accurate small LLM for data retrieval with a large LLM to deliver structured insights at enterprise scale reliably.

PROJECTS

Team Up London | TypeScript, React Native, Expo, Supabase, Google Maps API, Vercel   May 2025 – Jun 2025
- Built a full-stack mobile application using React Native and TypeScript with Supabase backend, enabling recent graduates and working professionals to discover social sports communities and impromptu games, serving 20+ active users.
- Established a complete CI/CD pipeline with Jest testing, Expo EAS Build for cross-platform app generation, and Vercel-hosted distribution platform, implementing agile methodology with iterative design-feedback cycles.

WACC (Compiler Project) | Scala   Jan 2025 – Mar 2025
- Developed a WACC compiler frontend, building a lexer and parser for syntax analysis, creating an abstract syntax tree representation, implementing a symbol table for semantic checks, and designing descriptive error reporting to aid debugging.
- Implemented the backend using TAC intermediate representation with ARM32/AArch64 support, architecture-specific dependencies, code optimisations (constant propagation, folding, control flow analysis), and a standard math library.

PintOS (Operating Systems Project) | C   Oct 2024 – Dec 2024
- Enhanced OS kernel functionality by implementing timer-based thread synchronisation, advanced priority scheduling with priority donation, and the BSD scheduler, ensuring efficient multitasking and thread management.
- Developed a virtual memory subsystem, including paging, frame management with second-chance eviction, and supplementary page tables, enabling support for user programs, memory-mapped files, and stack growth.

ARMv8 Emulator, Assembler, and Visualiser | C   May 2024 – Jun 2024
- Engineered a cycle-accurate ARMv8 emulator and two-pass assembler, implementing precise register, memory, and instruction management including robust error handling in compliance with ARMv8 specifications.
- Developed a SDL2-based GUI for real-time emulation visualisation, featuring drag-and-drop assembly parsing and dynamic rendering of CPU states, registers, memory maps, and ALU.

AI Research Agent | Python   Aug 2023 – Sep 2023
- Developed a AI research agent using LangChain, OpenAI's LLM, Serper API, BeautifulSoup4, and X API for automated search, web scraping, cited summaries, and content posting. Integrated MongoDB Atlas for geolocation tracking.
- Implemented a Streamlit web app enabling anyone to use the AI research agent, integrating geospatial visualisation to map worldwide usage using geolocation data stored in MongoDB Atlas. Analysed engagement patterns of 70+ global users.